STATE Readiness Checklist

A DIAGNOSTIC FOR PRODUCTION AI SYSTEMS

STATE BEATS INTELLIGENCE

The Meta Architect · simonparis.ca

Production AI systems fail not because **models are too weak** — but because they lack **proper state management, checkpointing, observability, and deterministic boundaries.** Score your system below. Find out where it will break before production does.

THE FIVE PILLARS — SCORE EACH 0, 1, OR 2

```
2 = yes, confidently
1 = partially / in progress
0 = no, or not sure
```

S **STRUCTURED** explicit state schemas

If this system crashes mid-run, can you look at the last saved state and know exactly where it stopped — without reading any logs or conversation history?

score: 0 | 1 | 2

T **TRACEABLE** every step logged

Can you pull a complete trace for any session from last week — every LLM call, every tool call, every input and output — right now, in under 5 minutes?

score: 0 | 1 | 2

A **AUDITABLE** governance-ready

If a regulator asks today what data your system used for a specific decision last month — can you answer in under 30 minutes?

score: 0 | 1 | 2

T **TOLERANT** resumable after failure

If this workflow fails at step 6 of 10 right now — does it resume from step 6, or restart from step 1?

score: 0 | 1 | 2

E **EXPLICIT** deterministic boundaries

For every LLM call in this system — what stops a wrong output from becoming a real-world action? If the answer is "nothing," that's your gap.

score: 0 | 1 | 2

MY TOTAL SCORE _____ / 10

WHAT YOUR SCORE MEANS — AND WHERE TO START

0–3	Critical Risk	One incident away from a rebuild. Start with S — define your state schema before anything else moves.
4–5	High Risk	Failures are happening — you can't see them yet. Add T logging to one workflow this week. Visibility before fixes.
6–7	Developing	Survives demos. A is where regulation finds you. Write a decision record for one automated output before you ship.
8–9	Prod-Ready	Shippable. Close the last gap. Run the Reboot Test below before you call it production-ready.
10	STATE ✓	Auditable, resumable, observable. Verify quarterly. New workflows reset the score.

THE REBOOT TEST — BEFORE YOU CALL IT PRODUCTION-READY

1. **Restart the service mid-run.** Does it resume from where it stopped?
2. **Inject a failure at step 6 of 10.** Does the error path clear the lock and report cleanly?
3. **Take a dependency offline for 2 minutes.** Does the system degrade gracefully or fail silently?
4. **Mock the LLM node with a worst-case output.** Does it reach a real-world action, or does the validation gate catch it?

Most teams score 4–6. The workshop walks through closing each gap on a real production system — live teardown, no slides.

→ simonparis.ca/workshop — next session: April 2026